

SNP-AV1: Stacked Neural Pruners for the AV1 Intra-Frame Partition Search

Pablo Rosa

*Video Technology Research Group (ViTech)
Federal University of Pelotas
Pelotas, Brazil
pablo.rosa@inf.ufpel.edu.br*

Daniel Palomino

*Video Technology Research Group (ViTech)
Federal University of Pelotas
Pelotas, Brazil
dpalomino@inf.ufpel.edu.br*

Marcelo Porto

*Video Technology Research Group (ViTech)
Federal University of Pelotas
Pelotas, Brazil
porto@inf.ufpel.edu.br*

Luciano Agostini

*Video Technology Research Group (ViTech)
Federal University of Pelotas
Pelotas, Brazil
agostini@inf.ufpel.edu.br*

Abstract—The intra-frame coding of the AOMedia Video 1 (AV1) codec relies on a recursive block partition search that evaluates ten candidate shapes at each node of the partition tree. At the highest-quality configuration this search is exhaustive by construction and dominates the intra-frame encoding time. This paper presents the Stacked Neural Pruners for AV1 (SNP-AV1), a fast decision algorithm for the intra-frame partition search applying two neural networks at two distinct stages of a single node. The first reads thirty-six attributes of causal context, whose extraction cost is negligible, and decides before any rate-distortion evaluation how much of the subtree deserves to be opened. The second adds the measured rate-distortion cost of the undivided candidate and decides whether the AB and 4-way shapes, which take one third of the local search time and which no earlier pruner targeted, are evaluated at all. The two are stacked because the candidates they remove are disjoint. Experimental results show that SNP-AV1 achieves an average time savings of 18.74% with a Bjontegaard delta bitrate (BD-BR) of 0.59%, reaching 32.16% and 1.42% under a stronger calibration, and the second pruner buys time about 3.4 times cheaper than loosening the thresholds of the first. To the best of the authors’ knowledge, this is the first solution in the literature to prune the extended partitions of the AV1 intra-frame partition search with a learned model.

Index Terms—Video coding, AV1, computational effort reduction, intra-frame block partitioning, machine learning, neural networks.

I. INTRODUCTION

Video is the dominant workload of the Internet, and the share of users watching online video every month has kept growing over the last years [1]. Alongside the standards defined by the international standardization bodies, such as the High Efficiency Video Coding (HEVC) [2] and the Versatile Video Coding (VVC) [3], developed by the Motion Picture Expert Group (ISO/IEC MPEG) and by the Video Coding Experts Group (ITU-T VCEG), the Alliance for Open Media (AOMedia) maintains the AOMedia Video 1 (AV1) [4], which follows the state-of-the-art flow of hybrid video encoders, splitting the process into intra-frame prediction, inter-frame

prediction, transforms, quantization, in-loop filters and entropy coding, and introduced new techniques in each of these steps [5], [6]. This coding efficiency comes at the price of computational effort: AV1 was shown to be superior to HEVC under certain conditions [7], but encoding videos of the same bitrate and resolution took twice as long [8], a problem aggravated during UHD 4K content processing.

One of the largest individual costs of AV1 intra-frame coding is the block partitioning decision. Every superblock is submitted to a recursive search that, at each node of the partition tree, evaluates the rate-distortion (RD) cost [9] of ten candidate shapes and keeps the cheapest one, descending recursively whenever the quad-tree split wins. The search is exhaustive by construction, since the encoder can only determine the best shape after evaluating the block under each candidate, and the reference software libaom [10] performs it in full at the highest-quality setting, `cpu-used=0`.

There are a few works in the literature targeting the computational effort of AV1, and none of them deals with this search. The works [11]–[13] act on the intra prediction mode decision loop, while [14], [15] are hardware solutions for the AV1 intra prediction, whose target is the circuit and not the encoder search, and [16] uses decision trees to determine whether all inter-prediction modes need to be evaluated. Learning-based partitioning solutions do exist for VVC [17]–[19] and can be adapted to the AV1 coding tools, but they were designed for the quad-tree with nested multi-type tree and evaluated on another encoder. None of these papers allow a direct and fair comparison, and none of them acts on the recursive partition search of the AV1 intra-frame coding.

This paper presents the Stacked Neural Pruners for AV1 (SNP-AV1), a fast decision algorithm confined to the intra-frame partition search, which leaves every other encoding stage untouched. The main strategy of this solution is to remove candidates progressively, as stronger evidence becomes available: a first network decides how much of the




Fig. 1. The ten partition shapes evaluated at a node, grouped into the four families used in this work and annotated with the share of local search time each family consumes.

partition tree deserves to be opened, from causal context whose extraction cost is negligible, and a second decides whether the most expensive family of candidates is worth evaluating, from the rate-distortion evidence the encoder has measured by then. The two are stacked because the sets of candidates they remove are disjoint, which is the condition under which the savings of two pruners accumulate instead of competing for the same time. Over the eight UHD 4K video sequences of the Common Test Conditions (CTC) Class A1, SNP-AV1 reaches an average time savings of 18.74% with a Bjøntegaard delta bitrate (BD-BR) of 0.586%.

II. BLOCK PARTITIONING IN AV1

The basic unit of the AV1 partitioning is the superblock, which has 128×128 or 64×64 pixels, and the structure that subdivides it is called partition tree [5]. AV1 allows ten partition types at each node [6], which Fig. 1 names and organizes into the four families used throughout this paper: the two structural decisions, PARTITION_NONE, hereafter NONE, which codes the node undivided, and PARTITION_SPLIT, the quad-tree split into four square sub-blocks; the two rectangular shapes; the four asymmetric AB shapes, each combining one rectangular half with two square sub-blocks; and the two 4-way shapes, of 4:1 aspect ratio. The AB and the 4-way families are jointly called the extended partitions.

The evaluation order inside a node is fixed and defines which information has already been paid at each point of the flow: NONE, the quad-tree split, the rectangular shapes and finally the extended partitions. A pruner acting before the first evaluation is blind to any rate-distortion number, and one acting after NONE has the measured rate, distortion and RD cost of the undivided candidate, at the price of never saving the cost of producing them.

The cost distribution in Fig. 1 was measured with the native partition timing instrumentation of libaom over 875,317 decision nodes of the 4K videos described in Section III-B, at `cpu-used=0`, taking as local work the sum of the nine non-recursive candidates of a node. It splits into three blocks of comparable size: NONE takes 30.1% of that time, the two rectangular shapes take 35.6% and the six extended ones take 34.3%. No single extended shape exceeds 7.22%, which is why the weight of the family as a whole went unnoticed, and the cost concentrates in the large nodes, reaching 50.0% at 32×32 and 51.0% at 64×64 pixels, which are the nodes a 4K frame produces in abundance.




Fig. 2. Insertion points of the two pruners in the search of a node, and the cascade of conditions each of them applies.

This work is focused on the acceleration of the recursive block partitioning decision of the AV1 intra-frame coding. The other decisions taken inside a coding block, such as the intra prediction modes, the transform type search and the in-loop filtering, lie outside the partition search loop attacked here, as does the inter-frame partition search, since the solution is trained and evaluated exclusively under the All-Intra (AI) encoder configuration.

III. PROPOSED SNP-AV1 SOLUTION

The Stacked Neural Pruners for AV1 (SNP-AV1) is a fast decision algorithm composed of two neural pruners inserted at two distinct stages of the search of a node, as shown in Fig. 2. The main strategy of this solution is to consult each pruner at the earliest point of the flow at which the information it needs has already been paid for.

The two stages are complementary rather than competing. The first acts before any rate-distortion evaluation, so a commitment removes every competing candidate of the node and the whole subtree below it, including the evaluation of NONE at each descendant; it pays for that reach by deciding blind to rate-distortion. The second acts only on the nodes that survive the first and reads the measured RD cost of NONE, over a narrower scope: the AB and 4-way candidates still standing.

The rule behind the pairing is that two pruners inserted in the same search compete for the same saved time whenever they remove the same candidates, however different the information each of them reads, and accumulate whenever the sets they remove are disjoint. The first stage asks whether the node can be settled at once, on the structural axis of NONE against the quad-tree split; the second asks whether the extended partitions, the block of cost isolated in Section II, are worth evaluating. One limit follows from the same rule, and Section IV measures it: a node the first stage has committed never reaches the second, so the disjunction that makes the two additive is a property of the operating point, and not of the pruners.

A. Attribute Design

The attribute vector of the first stage has thirty-six entries. Twenty-four are luma descriptors of the node and of its hierarchical context, covering variance, gradient energy and orientation, row and column profiles, edge strength and density, and the contrast against the parent and the sibling sub-blocks. Eight carry the causal partitioning neighborhood, that is, the availability of the above and left neighbors, their decided widths and heights in log2 units and their granularity and anisotropy; the last four carry the effective dequantization step of the luma DC coefficient, the normalized position in the frame and the node depth.

The criterion for inclusion was the cost of obtaining each attribute where the pruner is consulted: the twelve neighborhood, quantization and position entries are a free read of state the native search already maintains, and the luma descriptors require only arithmetic linear in the area of the node. The second stage reads these thirty-six plus the logarithm of the rate, of the distortion and of the RD cost measured for NONE. Every luma sample is brought to the eight-bit scale before any descriptor is computed, and the dequantization step with it, so no attribute depends on the bit depth and the models trained on eight-bit material are deployed unchanged on the ten-bit evaluation sequences.

B. Dataset Construction

The reference software libaom v3.10.0 was instrumented under a compile-time guard to record, at every node of the intra-frame partition search, the decision taken by the encoder itself together with the luma samples, the causal context, the quantization state, the position and the rate-distortion result of NONE. The guard is necessary because a binary that writes to disk at every node cannot measure wall-clock acceleration.

The dataset was built from the sixteen 4K sequences of the UVG dataset [20], at 3840×2160, 8-bit depth and 4:2:0 chroma subsampling, each encoded at CQ 20, 32, 43 and 55 over five frames sampled uniformly along the whole clip, since in AI coding consecutive frames are nearly identical. The extraction ran entirely at `cpu-used=0`, because the label is only faithful if it comes from a search that evaluated every candidate shape: an extraction at `cpu-used=6` produced zero counts for the six extended classes, exactly the ones the second stage must predict.

The resulting dataset contains 26.98 million samples, of which 10.07 million are decision nodes of 16, 32 and 64 pixels; the rest are 8×8 leaves whose label is constant and which are left out of the modeling. It was partitioned by sequence into ten training, three validation and three reserved test sequences, frozen before any measurement, and the eight evaluation sequences belong to the CTC Class A1 set, absent from the UVG corpus, so no sequence is reused across training, testing and evaluation, which rules out data leakage.

C. Network Architecture and Training

Both stages use multilayer perceptrons, one per node size, trained independently for the 16×16, 32×32 and 64×64 lev-

els; the 8×8 level is terminal and is not modeled. The topology is common to all six networks and deliberately shallow: two hidden layers of 64 and 32 rectified units and a linear output whose softmax gives three class probabilities in the first stage and two in the second, for 27,759 parameters in total. They were trained in PyTorch [21] with the AdamW optimizer [22], a learning rate of 1×10^{-3} , a weight decay of 1×10^{-4} , batches of 4096 samples and a fixed budget of 30 epochs, under hard-label cross entropy without class weighting. The attributes are standardized during training and the standardization is folded into the first linear layer at export time, so the encoder feeds raw attributes and the models execute natively within the encoding pipeline; this parity was verified over 192 random vectors, with a maximum difference of 1.35×10^{-7} .

On the reserved sequences, and before any encoder integration, the expected calibration error of the first stage is 0.0112 over 1,816,393 held-out decision nodes, so reading its thresholds as confidence levels is verified rather than assumed, and the three rate-distortion attributes raise the area under the receiver operating characteristic curve of the second from 0.890 to 0.902, enough to avoid half of the extended searches while losing 1.1% of their true winners.

D. Pruning Policies and Encoder Integration

The first-stage pruner applies three actions in cascade, drawn in Fig. 2, on the probabilities p_{none} , p_{split} and p_{rest} of the undivided, split and remaining classes. If $p_{\text{none}} > \tau_{\text{none}}$, the node is decided as NONE and the whole subtree is cut, which is the primary lever of time savings. Otherwise, if $p_{\text{split}} > \tau_{\text{split}}$, NONE is skipped as well and the search recurses only into the four sub-blocks. Otherwise, if $p_{\text{rest}} < \tau_{\text{rest}}$, the rectangular candidates are disabled, and the extended ones with them by dependency; otherwise the node proceeds to the full search. Two operating points were frozen before the final campaign: the balanced one uses $\tau_{\text{none}} = 0.95$, $\tau_{\text{split}} = 0.90$ and $\tau_{\text{rest}} = 0.20$, the aggressive one 0.60, 0.85 and 0.40.

The second-stage pruner marks the node to skip the extended partitions when p_{ext} , the estimated probability that an AB or a 4-way shape wins the node, falls below a threshold θ kept per node size, leaving NONE, the quad-tree split and the rectangular shapes untouched. The per-level threshold is experimental, since the values that discard the same fraction of true winners differ widely: calibrated on the reserved sequences to lose a tenth of the nodes whose optimal shape is extended, the deployed values are 0.0910, 0.1031 and 0.0144 at the 16, 32 and 64-pixel levels, and a stronger calibration gives 0.1776, 0.1627 and 0.0321.

Both pruners were implemented in C inside libaom v3.10.0, each behind a compile-time guard whose disabled build reproduces the reference bitstream byte for byte, which is what allows the marginal contribution of the second to be attributed to its decisions and not to a difference of build. The policy only removes candidates and never forces an illegal one, so SNP-AV1 is fully compliant with the AV1 specification and does not affect the bitstream syntax.

IV. RESULTS AND COMPARISONS

The developed solution was evaluated using the libaom v3.10.0 AV1 reference software, following the AOM Common Test Conditions (CTC) [23] and under the AI encoder configuration. The eight 4K sequences of the CTC Class A1 test set, named in Table I, are all at 3840×2160, 10 bits and 4:2:0 [24]. The first 15 frames of each were encoded, as the CTC specifies for the AI scenario, at CQ 20, 32, 43 and 55, against the unmodified libaom at `cpu-used=0` as anchor. The BD-BR [25] is computed over the luma peak signal-to-noise ratio, and all values reported here are positive, hence coding efficiency losses. Writing $t_{\text{cfg}}(s, q)$ and $t_{\text{anc}}(s, q)$ for the encoding times of the evaluated configuration and of the anchor on sequence s at quality point q , over the sequence set S and the quality set Q ,

$$\text{TS} = \frac{1}{|S|} \sum_{s \in S} \frac{1}{|Q|} \sum_{q \in Q} \left(1 - \frac{t_{\text{cfg}}(s, q)}{t_{\text{anc}}(s, q)} \right), \quad (1)$$

reported as a percentage. All BD-BR and TS results come from this grid; the shares in Fig. 1 were measured on the UVG corpus of Section III-B instead. Every configuration was encoded once, and repeating the anchor and the balanced first stage five times on Crosswalk gives a standard deviation of 0.23 percentage points on TS, so a difference of time savings below twice that value is not resolved. The BD-BR values carry no such caveat, since the byte count and the luma PSNR are deterministic.

One scope decision governs the comparisons that follow. The anchor is the exhaustive partition search itself, and no native speed preset is placed beside SNP-AV1: a preset reconfigures dozens of heuristics spread over every coding stage, whereas SNP-AV1 acts inside the partition search alone, so a side-by-side reading would rank arrangements of different scope rather than partition-search techniques.

Fig. 3 places the six measured operating points of SNP-AV1 against that anchor. The balanced calibration spans time savings of 17.72% to 19.81% at a BD-BR of 0.568% to 0.651%, and the aggressive one 31.51% to 32.16% at 1.403% to 1.420%. Inside each family the three points differ only in how much of the extended family the second pruner removes, and moving from one family to the other is a matter of thresholds alone, without retraining and without recompiling.

The second stage is measured marginally, because it acts only on the residue of nodes the first stage did not commit. Over the balanced base it contributes 1.02 percentage points of time savings, about 4.4 times the measured standard deviation, for 0.018 percentage points of BD-BR. That is a price of 0.018 point of coding efficiency per point of time bought, against the 0.061 charged by the only alternative the deployed solution leaves to its user, which is loosening the thresholds of the first stage: the second pruner is about 3.4 times cheaper, and the left window of Fig. 3 shows the two prices as the two slopes that leave the same operating point. Its stronger calibration costs 0.040 over the same base, still below the price of the threshold knob, which is why the milder one is deployed. Sequence by

TABLE I
PER-SEQUENCE RESULTS OF THE TWO DEPLOYED OPERATING POINTS
AGAINST THE LIBAOM V3.10.0 ANCHOR AT `cpu-used=0`

Sequence	Balanced		Aggressive	
	BD-BR	TS	BD-BR	TS
BoxingPractice	0.662%	20.33%	1.560%	31.97%
Crosswalk	0.534%	20.53%	1.145%	35.08%
FoodMarket2	0.609%	14.14%	1.386%	26.74%
Neon1224	0.685%	22.00%	1.427%	33.05%
NocturneDance	0.168%	13.78%	1.209%	30.79%
PierSeaSide	0.458%	17.47%	1.403%	26.03%
Tango	1.139%	22.89%	2.206%	38.80%
TimeLapse	0.433%	18.78%	0.934%	31.00%
Average	0.586%	18.74%	1.409%	31.68%




Fig. 3. Operating space of SNP-AV1 over the CTC Class A1 4K sequences, measured against the libaom v3.10.0 anchor at `cpu-used=0`, which is the origin of both axes. The time axis is broken into two windows of the same scale, the balanced calibration on the left and the aggressive one on the right, over a common BD-BR axis. Inside each window the three points differ only in the strength of the second pruner; the dashed line is the threshold knob of the first stage, and the two annotated slopes are the prices of the two levers, in BD-BR points per point of time savings.

sequence, at time savings matched by interpolating the first-stage thresholds, the stacked point is the better of the two in six of the eight sequences, two of them by strict Pareto dominance, and it loses only in Crosswalk and Neon1224, the two sequences in which the extended axis is barely exercised and the second pruner has almost nothing to remove.

This additivity is a property of the operating point, and not of the pruner. Over the aggressive base the same second stage adds 0.17 percentage points, below the resolution of the arrangement and above it in one sequence out of eight instead of six. With the undivided threshold loosened to 0.60, the first stage commits the nodes so early that they seldom reach the point at which the extended partitions would be decided, and the two pruners fall back to disputing the same residue. The right window of Fig. 3 is that collapse: at the same scale as the left one, the whole aggressive family spans 0.65 percentage points of time savings, against the 2.09 the second pruner alone covers over the balanced base.

Table I opens the two deployed points per sequence, and

what it exposes is content dispersion, not a swept range. Read as time bought per point of coding efficiency ceded, the median of the eight sequences is 35 points of time savings per point of BD-BR at the balanced point and 22 at the aggressive one, the same rising price the aggregate already showed. Tango is the least favorable sequence under both calibrations, and the two columns move together across the set: the sequences in which the first stage commits the most nodes are also the ones in which committing costs the most.

The overhead is reported rather than omitted: the complete path of SNP-AV1, attribute extraction and inference together, stays below one third of one percent of the encoding time, so the time savings come from the pruning decisions and not from the lightness of each inference.

V. CONCLUSIONS

This paper presented SNP-AV1, a machine-learning solution that stacks two neural pruners at two distinct stages of the AV1 intra-frame partition search, reaching an average time savings of 18.74% with a BD-BR of 0.59%, and 32.16% with 1.42% under a stronger calibration. The first pruner decides how much of the subtree is opened before any rate-distortion cost is paid, and the second whether the extended partitions, one third of the local search time, are worth evaluating; the two are stacked because the candidates they remove are disjoint. Both were implemented in C inside libaom v3.10.0 and evaluated on the eight 4K sequences of the CTC Class A1 under the AI configuration, against the exhaustive search as anchor. Measured marginally over the first, the second pruner adds 1.02 percentage points of time savings for 0.018 of BD-BR, about 3.4 times cheaper than buying the same time from the thresholds of the first stage, and it is the better of the two in six of the eight sequences; the advantage collapses on the aggressive base, where the first stage commits the nodes before the second is consulted. Besides, to the best of the authors' knowledge, this is the first solution in the literature to prune the extended partitions of the AV1 intra-frame partition search with a learned model.

ACKNOWLEDGMENT

This study was financed in part by the Coordenação de Aperfeiçoamento de Pessoal de Nível Superior – Brasil (CAPES) – Finance Code 001, and also by the Brazilian research support agencies CNPq and FAPERGS.

REFERENCES

- [1] Statista. (2025). *Share of Internet Users Watching Online Videos Every Month From 1st Quarter 2022 to 3rd Quarter 2024*. [Online]. Available: <https://www.statista.com/statistics/1489445/online-video-viewers-worldwide-quarterly/>
- [2] ITU-T. (2013). *H.265: High Efficiency Video Coding*. [Online]. Available: <https://www.itu.int/rec/t-rec-h.265>
- [3] ITU-T. (2020). *H.266: Versatile Video Coding*. [Online]. Available: <https://www.itu.int/rec/t-rec-h.266>
- [4] Alliance for Open Media. (2019). *AV1 Video Codec*. [Online]. Available: <https://aomedia.org/av1>
- [5] J. Han *et al.*, "A technical overview of AV1," *Proc. IEEE*, vol. 109, no. 9, pp. 1435–1462, Sep. 2021.
- [6] Y. Chen *et al.*, "An overview of core coding tools in the AV1 video codec," in *Proc. Picture Coding Symp. (PCS)*, Jun. 2018, pp. 41–45.
- [7] D. Grois, T. Nguyen, and D. Marpe, "Coding efficiency comparison of AV1/VP9, H.265/MPEG-HEVC, and H.264/MPEG-AVC encoders," in *Proc. Picture Coding Symp. (PCS)*, Dec. 2016, pp. 1–5.
- [8] M. A. Layek *et al.*, "Performance analysis of H.264, H.265, VP9 and AV1 video encoders," in *Proc. 19th Asia-Pacific Netw. Oper. Manage. Symp. (APNOMS)*, Sep. 2017, pp. 322–325.
- [9] G. J. Sullivan and T. Wiegand, "Rate-distortion optimization for video compression," *IEEE Signal Process. Mag.*, vol. 15, no. 6, pp. 74–90, Nov. 1998.
- [10] Alliance for Open Media. *AV1 Codec Library*, v3.10.0. [Online]. Available: <https://aomedia.googlesource.com/aom>
- [11] M. Xu and B. Jeon, "Selection of intra prediction tools for fast AV1 encoding," in *Proc. IEEE Int. Symp. Broadband Multimedia Syst. Broadcast. (BMSB)*, Oct. 2020, pp. 1–4.
- [12] M. Xu and B. Jeon, "User-priority based AV1 coding tool selection," *IEEE Trans. Broadcast.*, vol. 67, no. 3, pp. 736–745, Sep. 2021.
- [13] J. Jeong, G. Gankhuyag, and Y.-H. Kim, "A fast intra mode decision based on accuracy of rate distortion model for AV1 intra encoding," in *Proc. 34th Int. Tech. Conf. Circuits/Syst., Comput. Commun. (ITC-CSCC)*, Jun. 2019, pp. 1–3.
- [14] M. Corrêa, B. Zatt, D. Palomino, G. Corrêa, and L. Agostini, "A high-throughput hardware architecture for AV1 non-directional intra modes," *IEEE Trans. Circuits Syst. I, Reg. Papers*, vol. 67, no. 5, pp. 1481–1494, May 2020.
- [15] L. Netto, M. Corrêa, D. Palomino, L. Agostini, and G. Corrêa, "Power-quality configurable hardware design for AV1 directional intra-frame prediction," *IEEE Design Test*, vol. 39, no. 2, pp. 38–45, Apr. 2022.
- [16] J. Kim, S. Blasi, A. S. Dias, M. Mrak, and E. Izquierdo, "Fast inter-prediction based on decision trees for AV1 encoding," in *Proc. IEEE Int. Conf. Acoust., Speech Signal Process. (ICASSP)*, May 2019, pp. 1627–1631.
- [17] M. Saldanha, G. Sanchez, C. Marcon, and L. Agostini, "Configurable fast block partitioning for VVC intra coding using light gradient boosting machine," *IEEE Trans. Circuits Syst. Video Technol.*, vol. 32, no. 6, pp. 3947–3960, Jun. 2022.
- [18] Q. He, W. Wu, L. Luo, C. Zhu, and H. Guo, "Random forest based fast CU partition for VVC intra coding," in *Proc. IEEE Int. Symp. Broadband Multimedia Syst. Broadcast. (BMSB)*, Aug. 2021, pp. 1–4.
- [19] J. Chen, H. Sun, J. Katto, X. Zeng, and Y. Fan, "Fast QTMT partition decision algorithm in VVC intra coding based on variance and gradient," in *Proc. IEEE Vis. Commun. Image Process. (VCIP)*, Dec. 2019, pp. 1–4.
- [20] A. Mercat, M. Viitanen, and J. Vanné, "UVG dataset: 50/120fps 4K sequences for video codec analysis and development," in *Proc. 11th ACM Multimedia Syst. Conf. (MMSys)*, Jun. 2020, pp. 297–302.
- [21] A. Paszke *et al.*, "PyTorch: An imperative style, high-performance deep learning library," in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, Dec. 2019, pp. 8024–8035.
- [22] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *Proc. Int. Conf. Learn. Represent. (ICLR)*, May 2019, pp. 1–18.
- [23] X. Zhao *et al.*, "AOM common test conditions v9.0," Alliance for Open Media, Codec Working Group, doc. CWG-G082, Mar. 2026.
- [24] Xiph.Org Foundation. *AOM CTC Test Set, Class A1 (4K)*. [Online]. Available: https://media.xiph.org/video/aomctc/test_set/a1_4k/
- [25] G. Bjøntegaard, "Calculation of average PSNR differences between RD-curves," Video Coding Experts Group (VCEG), doc. VCEG-M33, Mar. 2001.